

CSCI 4140: Natural Language Processing

CSCI/DASC 6040: Computational Analysis of Natural Languages

Spring 2024

Homework 1 - N-gram models

Due Sunday, January 21, at 11:59 PM

Do not redistribute without the instructor's written permission.

The learning goals of this assignment are to:

- Understand how to compute language model probabilities using maximum likelihood estimation.
- Implement back-off.
- Have fun using a language model to probabilistically generate texts.
- Compare word-level language models and character-level language models.

✓ Part 1: N-gram Language model (60 pts)

✓ Preliminaries

```
import random
from collections import *
import numpy as np
```

We'll start by loading the data. The WikiText language modeling dataset is a collection of tokens extracted from the set of verified Good and Featured articles on Wikipedia.

```
data = {'test': '', 'train': '', 'valid': ''}

for data_split in data:
    fname = "wiki.{}.tokens".format(data_split)
    with open(fname, 'r') as f_wiki:
        data[data_split] = f_wiki.read().lower().split()

vocab = list(set(data['train']))
```

Now have a look at the data by running this cell.

```
print('train : %s ...' % data['train'][:10])
print('dev : %s ...' % data['valid'][:10])
print('test : %s ...' % data['test'][:10])
print('first 10 words in vocab: %s' % vocab[:10])

train : ['=', 'valkyria', 'chronicles', 'iii', '=', 'senjō', 'no', 'valkyria', '3', ':'] ...
dev : ['=', 'homarus', 'gamarus', '=', 'homarus', 'gamarus', ',', 'known', 'as', 'the'] ...
test : ['=', 'robert', '<unk>', '=', 'robert', '<unk>', 'is', 'an', 'english', 'film'] ...
first 10 words in vocab: ['noble', 'effigies', 'kovacs', '272', 'balista', '攻殻機動隊', '2e', 'creeks', 'laverton', 'refute']
```

✓ Q1.1: Train N-gram language model (25 pts)

Complete the following `train_ngram_lm` function based on the following input/output specifications. If you've done it right, you should pass the tests in the cell below.

Input:

- **data**: the data object created in the cell above that holds the tokenized Wikitext data
- **order**: the order of the model (i.e., the "n" in "n-gram" model). If `order=3`, we compute $p(w_2|w_0, w_1)$.

Output:

- **lm**: A dictionary where the key is the history and the value is a probability distribution over the next word computed using the maximum likelihood estimate from the training data. Importantly, this dictionary should include *backoff* probabilities as well; e.g., for `order=4`, we want to store $p(w_3|w_0, w_1, w_2)$ as well as $p(w_3|w_1, w_2)$ and $p(w_3|w_2)$.

Each key should be a single string where the words that form the history have been concatenated using spaces. Given a key, its corresponding value should be a dictionary where each word type in the vocabulary is associated with its probability of appearing after the key. For example,

the entry for the history 'w1 w2' should look like:

```
lm['w1 w2'] = {'w0': 0.001, 'w1' : 1e-6, 'w2' : 1e-6, 'w3': 0.003, ...}
```

In this example, we also want to store `lm['w2']` and `lm['']`, which contain the bigram and unigram distributions respectively.

Hint: You might find the **defaultdict** and **Counter** classes in the **collections** module to be helpful.

```
def train_ngram_lm(data, order=3):
    """
    Train n-gram language model
    """

    # pad (order-1) special tokens to the left
    # for the first token in the text
    order -= 1
    data = ['<S>'] * order + data
    lm = defaultdict(Counter)
    for i in range(len(data) - order):
        history = ' '.join(data[i:i+order])
        current_word = data[i+order]

        # Update language model
        for j in range(order):
            partial_history = ' '.join(data[i+j:i+order])
            lm[partial_history][current_word] += 1

        # Update unigram model
        lm[''][current_word] += 1

    # Convert counts to probabilities
    for history, word_counts in lm.items():
        total_count = sum(word_counts.values())
        probabilities = {word: count / total_count for word, count in word_counts.items()}
        lm[history] = probabilities

    return lm

def test_ngram_lm():

    print('checking empty history ...')
    lm1 = train_ngram_lm(data['train'], order=1)
    assert '' in lm1, "empty history should be in the language model!"

    print('checking probability distributions ...')
    lm2 = train_ngram_lm(data['train'], order=2)
    sample = [sum(lm2[k].values()) for k in random.sample(list(lm2), 10)]
    assert all([a > 0.999 and a < 1.001 for a in sample]), "lm[history][word] should sum to 1!"

    print('checking lengths of histories ...')
    lm3 = train_ngram_lm(data['train'], order=3)
    assert len(set([len(k.split()) for k in list(lm3)])) == 3, "lm object should store histories of all sizes!"

    print('checking word distribution values ...')
    assert lm1['']['the'] < 0.064 and lm1['']['the'] > 0.062 and \
        lm2['the']['first'] < 0.017 and lm2['the']['first'] > 0.016 and \
        lm3['the first']['time'] < 0.106 and lm3['the first']['time'] > 0.105, \
        "values do not match!"

    print("Congratulations, you passed the ngram check!")

test_ngram_lm()

checking empty history ...
checking probability distributions ...
checking lengths of histories ...
checking word distribution values ...
Congratulations, you passed the ngram check!
```

✓ Q1.2: Generate text from n-gram language model (10 pts)

Complete the following `generate_text` function based on these input/output requirements:

Input:

- **lm:** the `lm` object is the dictionary you return from the `train_ngram_lm` function
- **vocab:** `vocab` is a list of unique word types in the training set, already computed for you during data loading.
- **context:** the input context string that you want to condition your language model on, should be a space-separated string of tokens
- **order:** order of your language model (i.e., "n" in the "n-gram" model)
- **num_tok:** number of tokens to be generated following the input context

Output:

- generated text, should be a space-separated string

Hint:

After getting the next-word distribution given history, try using [numpy.random.choice](#) to sample the next word from the distribution.

```
# generate text
def generate_text(lm, vocab, context="he is the", order=3, num_tok=25):

    # The goal is to generate new words following the context
    # If context has more tokens than the order of lm,
    # generate text that follows the last (order-1) tokens of the context
    # and store it in the variable 'history'
    order -= 1
    history = context.split()[-order:]
    # `out` is the list of tokens of context
    # you need to append the generated tokens to this list
    out = context.split()

    for i in range(num_tok):
        # Get the history as a string
        history_str = ' '.join(history)

        # Check if the history is in the language model
        if history_str in lm:
            # Get the next-word distribution given the history
            next_word_dist = lm[history_str]

            # Sample the next word using numpy.random.choice
            next_word = np.random.choice(list(next_word_dist.keys()), p=list(next_word_dist.values()))

            # Append the next word to the output
            out.append(next_word)

            # Update the history for the next iteration
            history = out[-order:]
        else:
            # If history is not in the language model, break the loop
            break

    return ' '.join(out)
```

Now try to generate some texts, generated by ngram language model with different orders.

```
order = 1
generate_text(train_ngram_lm(data['train'], order=order), vocab, context='he is the', order=order)

'he is the'

order = 2
generate_text(train_ngram_lm(data['train'], order=order), vocab, context='he is the', order=order)

'he is the headshrinkers match around 5 million households . in wales : what if they attempted to
moments hide their first weekend but if he then lost'

order = 3
generate_text(train_ngram_lm(data['train'], order=order), vocab, context='he is the', order=order)

'he is the main factors preventing the transmission of information of the most important divinit
as of the cedarite arabs " heat of song in choreographed moves backed'

order = 4
generate_text(train_ngram_lm(data['train'], order=order), vocab, context='he is the', order=order)

'he is the only god , creator of the armenian alphabet . it was the first organelle to be discove
red and two younger children ( one around'
```

✓ Q1.3 : Evaluate the models (25 pts)

Now let's evaluate the models quantitatively using the intrinsic metric **perplexity**.

Recall perplexity is the inverse probability of the test text

$$PP(w_1, \dots, w_t) = P(w_1, \dots, w_t)^{-\frac{1}{t}}$$

For an n-gram model, perplexity is computed by

$$PP(w_1, \dots, w_t) = \left[\prod_{t=1}^T P(w_t | w_{t-1}, \dots, w_{t-n+1}) \right]^{-\frac{1}{T}}$$

To address the numerical issue (underflow), we usually compute

$$PP(w_1, \dots, w_t) = \exp\left(-\frac{1}{T} \sum_i \log P(w_i | w_{t-1}, \dots, w_{t-n+1})\right)$$

Input:

- **lm**: the language model you trained (the object you returned from the `train_ngram_lm` function)
- **data**: test data
- **vocab**: the list of unique word types in the training set
- **order**: order of the lm

Output:

- the perplexity of test data

Hint:

- If the history is not in the **lm** object, back-off to (n-1) order history to check if it is in **lm**. If no history can be found, just use $1/|V|$ where $|V|$ is the size of vocabulary.

```
from collections import Counter, defaultdict
from math import log, exp

def compute_perplexity(lm, data, vocab, order=3):
    # pad according to order
    order -= 1
    data = ['<S>'] * order + data
    log_sum = 0
    V = len(vocab)

    for i in range(len(data) - order):
        h, w = ' '.join(data[i: i+order]), data[i+order]

        # If history h is not in lm, back-off to (n-1) gram and look up again
        while h not in lm and order > 1:
            order -= 1
            h = ' '.join(data[i: i+order])

        # If no history can be found, use 1/|V|
        if h not in lm:
            log_sum += -log(1/V)
        else:
            # Look up probability in the language model
            log_sum += -log(lm[h].get(w, 1/V))

    # Compute perplexity
    perplexity = exp(log_sum / len(data))
    return perplexity
```

Let's evaluate the language model with different orders. You should see a decrease in perplexity as the order increases. As a reference, the perplexity of the unigram, bigram, trigram, and 4-gram language models should be around 795, 203, 141, and 130 respectively.

```
for o in [1, 2, 3, 4]:
    lm = train_ngram_lm(data['train'], order=o)
    print('order {} ppl {}'.format(o, compute_perplexity(lm, data['test'], vocab, order=o)))

order 1 ppl 652.7888243015223
order 2 ppl 289.18308911565356
order 3 ppl 289.21504974686553
order 4 ppl 289.2197909653991
```

✓ Part 2: Character-level N-gram language model (50 points)

In the lecture, language modeling was defined as the task of predicting the next word in a sequence given the previous words. In this part of the assignment, we will focus on the related problem of predicting the next character or word in a sequence given the previous characters.

✓ Preliminaries

We have to modify how we load the data, by splitting it into characters rather than words. Also, we'll use both upper case and lower case letters.

```
data = {'test': '', 'train': '', 'valid': ''}

for data_split in data:
    fname = "wiki.{}.tokens".format(data_split)
    data[data_split] = list(open(fname, 'r', encoding = 'utf-8').read())

vocab = list(set(data['train']))
```

Now have a look at the data by running this cell.

```
print('train : %s ...' % data['train'][:10])
print('dev : %s ...' % data['valid'][:10])
print('test : %s ...' % data['test'][:10])
print('first 10 characters in vocab: %s' % vocab[:10])

train : [' ', '\n', ' ', '=', ' ', 'V', 'a', 'l', 'k', 'y'] ...
dev : [' ', '\n', ' ', '=', ' ', 'H', 'o', 'm', 'a', 'r'] ...
test : [' ', '\n', ' ', '=', ' ', 'R', 'o', 'b', 'e', 'r'] ...
first 10 characters in vocab: ['ê', 'ġ', '±', '-', '£', 'ä', '.', 'X', 'ç', 'Ħ']
```

✓ Q2.1: Train N-gram language model (20 pts)

Complete the following `train_ngram_lm` function based on the following input/output specifications. If you've done it right, you should pass the tests in the cell below.

Input:

- **data**: the data object created in the cell above that holds the tokenized Wikitext data
- **order**: the order of the model (i.e., the "n" in "n-gram" model). If `order=3`, we compute $p(c_2 | c_0, c_1)$.

Output:

- **lm**: A dictionary where the key is the history and the value is a probability distribution over the next character computed using the maximum likelihood estimate from the training data. Importantly, this dictionary should include *backoff* probabilities as well; e.g., for `order=4`, we want to store $p(c_3 | c_0, c_1, c_2)$ as well as $p(c_3 | c_1, c_2)$ and $p(c_3 | c_2)$.

Each key should be a single string where the characters that form the history have been concatenated. Given a key, its corresponding value should be a dictionary where each character in the vocabulary is associated with its probability of appearing after the key. For example, the entry for the history 'c1c2' should look like:

```
lm['c1c2'] = {'c0': 0.001, 'c1': 1e-6, 'c2': 1e-6, 'c3': 0.003, ...}
```

In this example, we also want to store `lm['c2']` and `lm['']`, which contain the bigram and unigram distributions respectively.

Hint: You might find the **defaultdict** and **Counter** classes in the **collections** module to be helpful.

```
def train_ngram_lm(data, order=3):
    """
    Train n-gram language model
    """

    # pad (order-1) special tokens to the left
    # for the first token in the text
    order -= 1
    data = ['~'] * order + data #
    lm = defaultdict(Counter)
    for i in range(len(data) - order):
        history = ''.join(data[i: i+order])
        lm[history][data[i+order]] += 1

    # Convert counts to probabilities
    for history, counts in lm.items():
        total_count = sum(counts.values())
        lm[history] = {char: count / total_count for char, count in counts.items()}

    return lm
```

```
def test_ngram_lm():

    print('checking empty history ...')
    lm1 = train_ngram_lm(data['train'], order=1)
    assert ' ' in lm1, "empty history should be in the language model!"

    print('checking probability distributions ...')
    lm2 = train_ngram_lm(data['train'], order=2)
    sample = [sum(lm2[k].values()) for k in random.sample(list(lm2), 10)]
    assert all([a > 0.999 and a < 1.001 for a in sample]), "lm[history][character] should sum to 1!"

    print('checking lengths of histories ...')
    lm3 = train_ngram_lm(data['train'], order=3)
    #assert len(set([len(k) for k in list(lm3)])) == 3, "lm object should store histories of all sizes!"

    print('checking character distribution values ...')
    assert lm1[' ']['t'] < 0.062 and lm1[' ']['t'] > 0.060 and \
        lm2['t']['h'] < 0.297 and lm2['t']['h'] > 0.296 and \
        lm3['th']['e'] < 0.694 and lm3['th']['e'] > 0.693, \
        "values do not match!"

    print("Congratulations, you passed the ngram check!")

test_ngram_lm()

    checking empty history ...
    checking probability distributions ...
    checking lengths of histories ...
    checking character distribution values ...
    Congratulations, you passed the ngram check!
```

✓ Q2.2: Generate text from n-gram language model (10 pts)

Complete the following `generate_text` function based on these input/output requirements:

Input:

- **lm**: the `lm` object is the dictionary you return from the `train_ngram_lm` function
- **vocab**: `vocab` is a list of unique characters in the training set, already computed for you during data loading.
- **context**: the input context string that you want to condition your language model on, should be a string
- **order**: order of your language model (i.e., "n" in the "n-gram" model)
- **num_tok**: number of characters to be generated following the input context

Output:

- generated text, should be a sequence of characters

Hint:

After getting the next-character distribution given history, try using [numpy.random.choice](#) to sample the next character from the distribution.

```
# generate text
def generate_text(lm, vocab, context="he ", order=3, num_tok=25):

    # The goal is to generate new characters following the context
    # If context has more tokens than the order of lm,
    # generate text that follows the last (order-1) tokens of the context
    # and store it in the variable `history`
    order -= 1
    history = list(context)[-order:]
    # `out` is the list of tokens of context
    # you need to append the generated tokens to this list
    out = list(context)

    for i in range(num_tok):
        # Get the history as a string
        history_str = ''.join(history)

        # Check if the history is in the language model
        if history_str in lm:
            # Get the next-character distribution given the history
            next_char_dist = lm[history_str]

            # Sample the next character using numpy.random.choice
            next_char = np.random.choice(list(next_char_dist.keys()), p=list(next_char_dist.values()))

            # Append the next character to the output
            out.append(next_char)

            # Update the history for the next iteration
            history = out[-order:]
        else:
            # If history is not in the language model, break the loop
            break

    return ''.join(out)
```

Now try to generate some texts, generated by ngram language model with different orders.

```
order = 1
generate_text(train_ngram_lm(data['train'], order=order), vocab, context='he is the', order=order)

'he is the'

order = 2
generate_text(train_ngram_lm(data['train'], order=order), vocab, context='he is the', order=order)

'he is thed sth be lernieas 20 ) ct'

order = 3
generate_text(train_ngram_lm(data['train'], order=order), vocab, context='he is the', order=order)

'he is thernits the fire Val bee fa'

order = 4
generate_text(train_ngram_lm(data['train'], order=order), vocab, context='he is the', order=order)

'he is the Piazz , Prespecial bes ,'
```

✓ Q2.3 : Evaluate the models (20 pts)

Now let's evaluate the models quantitatively using the intrinsic metric **perplexity**.

Recall perplexity is the inverse probability of the test text

$$PP(w_1, \dots, w_t) = P(w_1, \dots, w_t)^{-\frac{1}{T}}$$

For an n-gram model, perplexity is computed by

$$PP(w_1, \dots, w_t) = \left[\prod_{t=1}^T P(w_t | w_{t-1}, \dots, w_{t-n+1}) \right]^{-\frac{1}{T}}$$

To address the numerical issue (underflow), we usually compute

$$PP(w_1, \dots, w_t) = \exp \left(-\frac{1}{T} \sum_i \log P(w_t | w_{t-1}, \dots, w_{t-n+1}) \right)$$

Input:

- **lm**: the language model you trained (the object you returned from the `train_ngram_lm` function)
- **data**: test data

- **vocab**: the list of unique characters in the training set
- **order**: order of the lm

Output:

- the perplexity of test data

Hint:

- If the history is not in the **lm** object, back-off to (n-1) order history to check if it is in **lm**. If no history can be found, just use $1/|V|$ where $|V|$ is the size of vocabulary.

```
from math import log, exp
def compute_perplexity(lm, data, vocab, order=3):

    # pad according to order
    order -= 1
```